

Rasco: Resource Allocation and Scheduling Co-design for DAG Applications on Multicore

Abigail Eisenklam, Robert Gifford, Georgiy A. Bondar*, Yifan Cai, Tushar Sial^T,
Linh Thi Xuan Phan, Abhishek Halder^{T*}



UC SANTA CRUZ*

IOWA STATE^T
UNIVERSITY

The Rise of Data Intensive CPS Tasks

- Tasks in real-time, embedded, and CPS are increasingly **data intensive**
- Autonomous control, image processing, signal processing, etc.



Joby Acquires Xwing Autonomy Division, Looks Ahead to Autonomous Flight

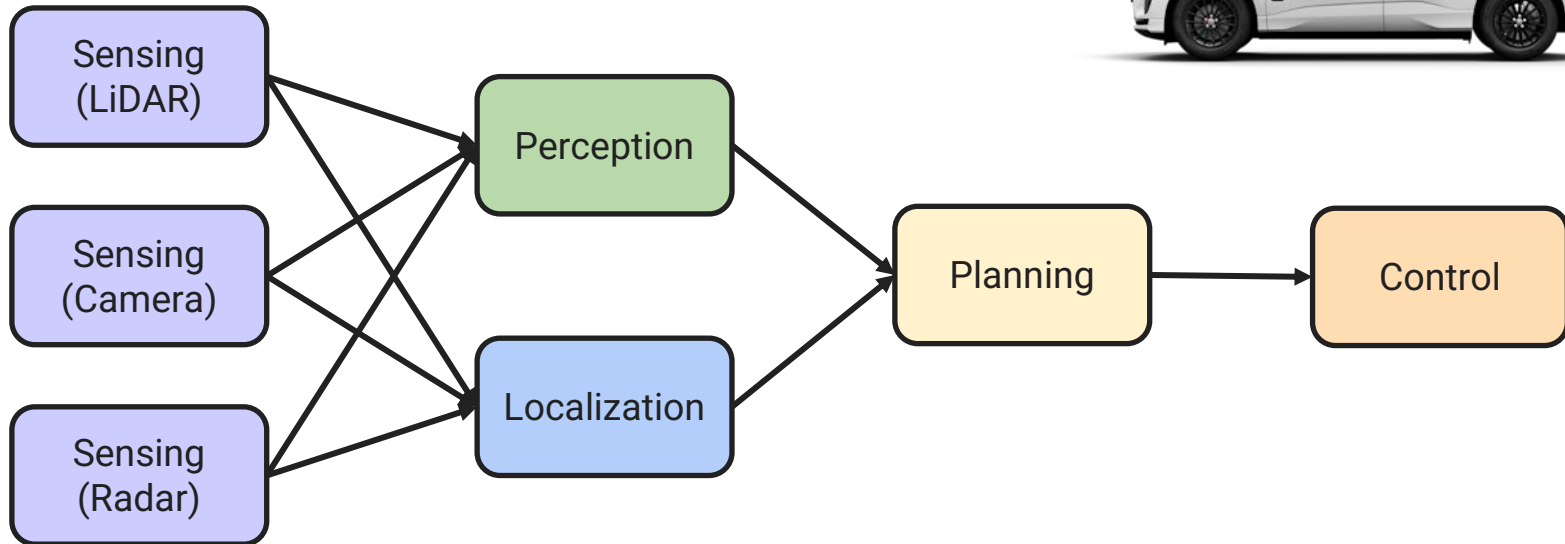
June 04, 2024 7:00am EDT

 [Download as PDF](#)

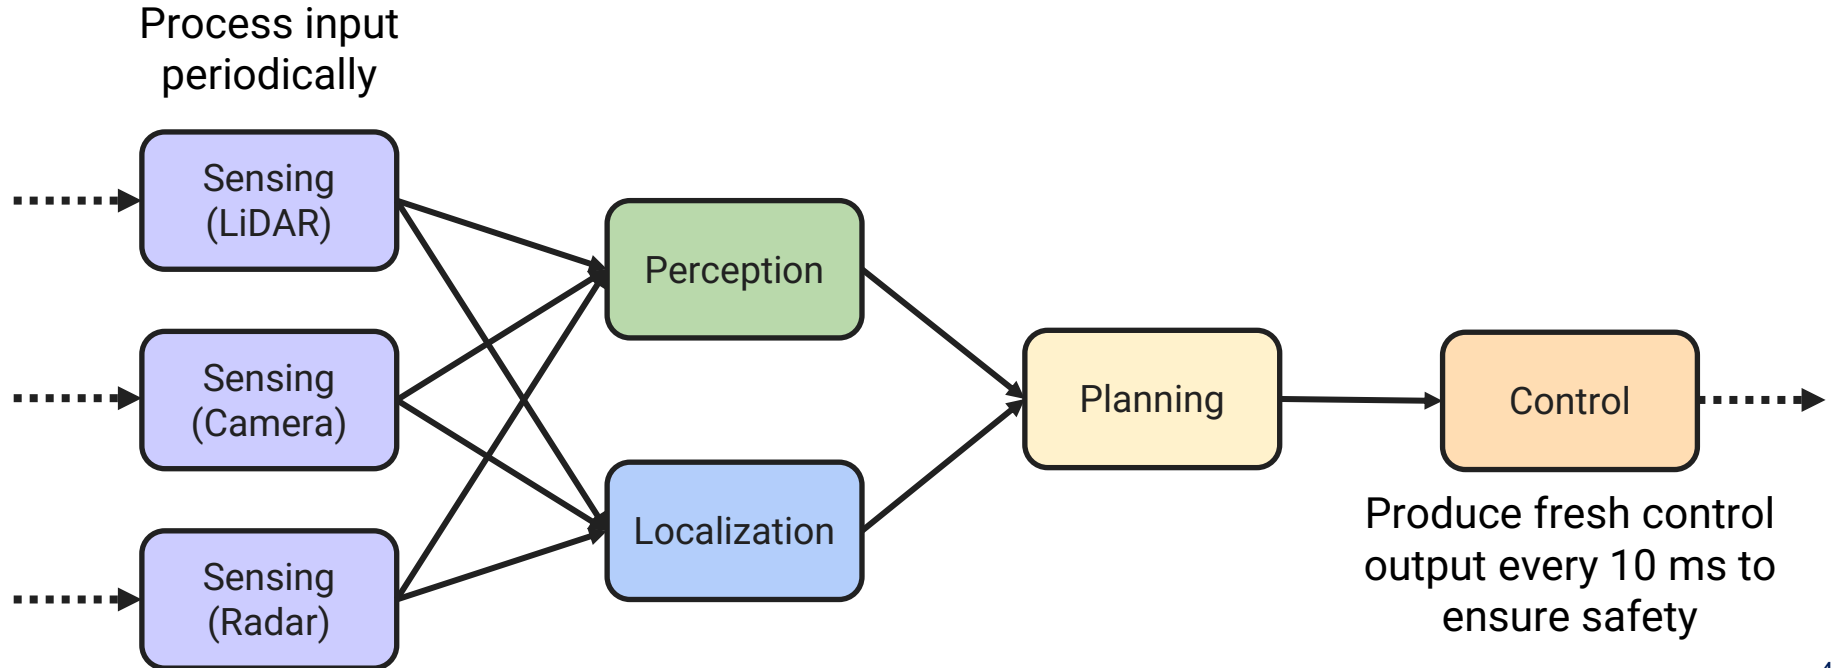
<https://www.jobyaviation.com/>

Data Dependencies in CPS

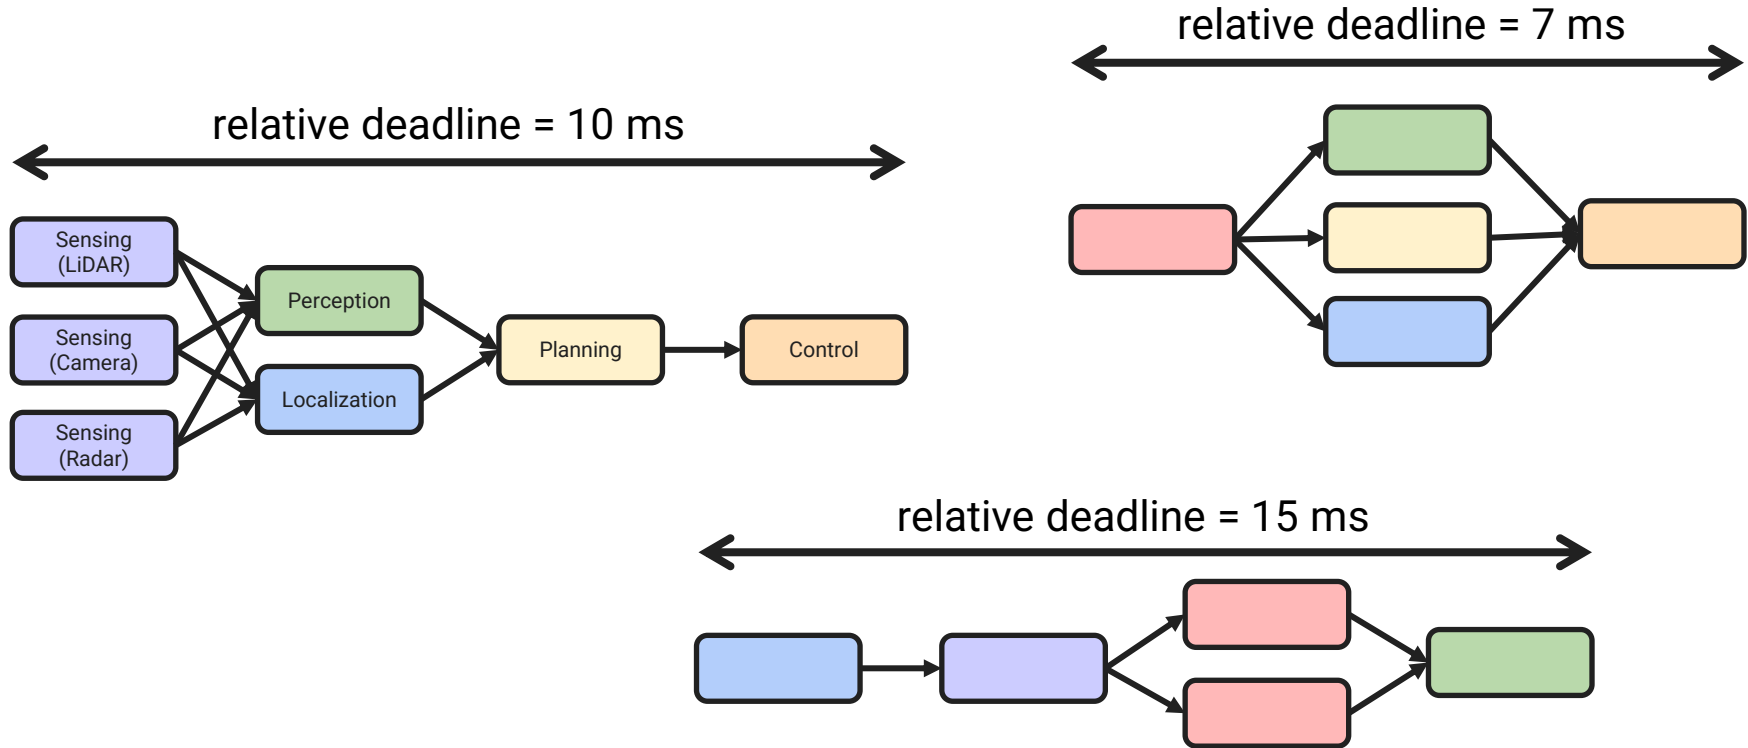
- CPS tasks have data dependencies
- Example: [Autoware](#) pipeline



Data Dependencies in CPS

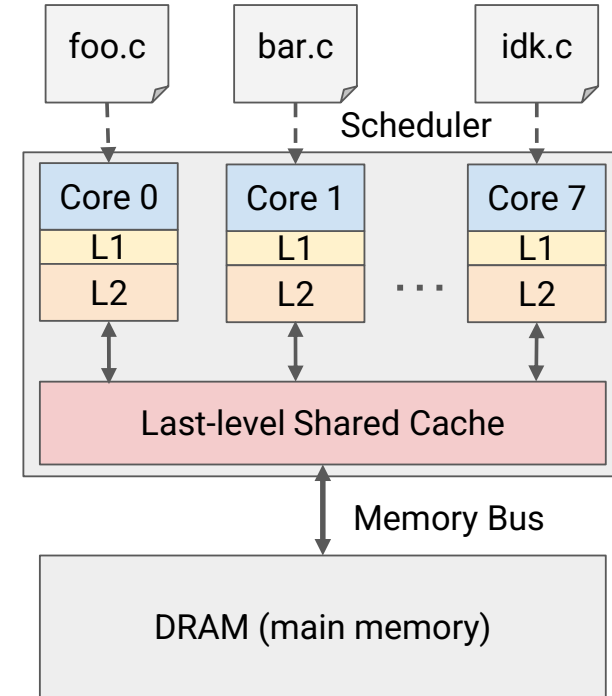


Model: Multiple Periodic DAG Tasks w/ Implicit Deadlines




Leveraging Multicore Hardware in CPS

- Multicores exploit inter- and intra-DAG parallelism
 - Lower DAG latency
 - Higher system throughput




Leveraging Multicore Hardware in CPS

- Multicores exploit inter- and intra-DAG parallelism
 - Lower DAG latency
 - Higher system throughput
- Shared resources such as last-level cache and memory bandwidth are statistically multiplexed
 - Good average case performance



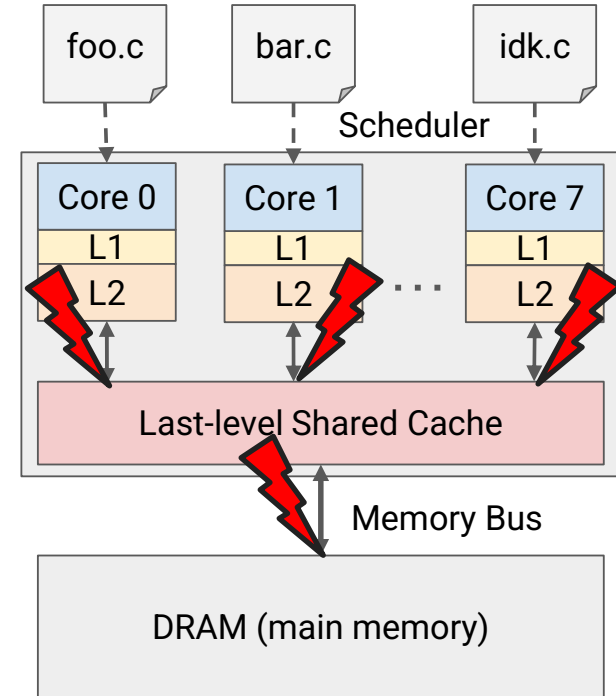
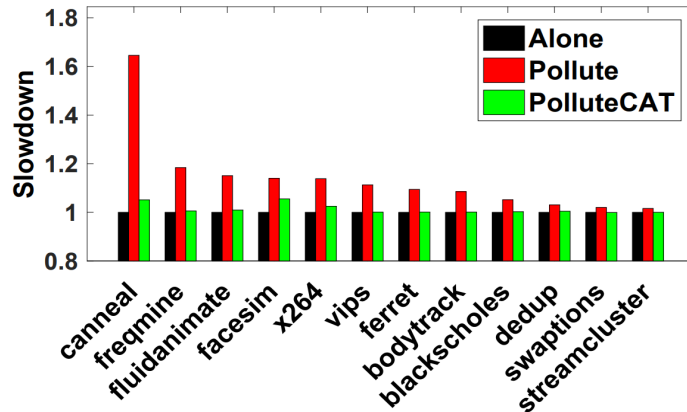
Challenges on Multicore

- Contention for shared resources can cause **interference** between tasks



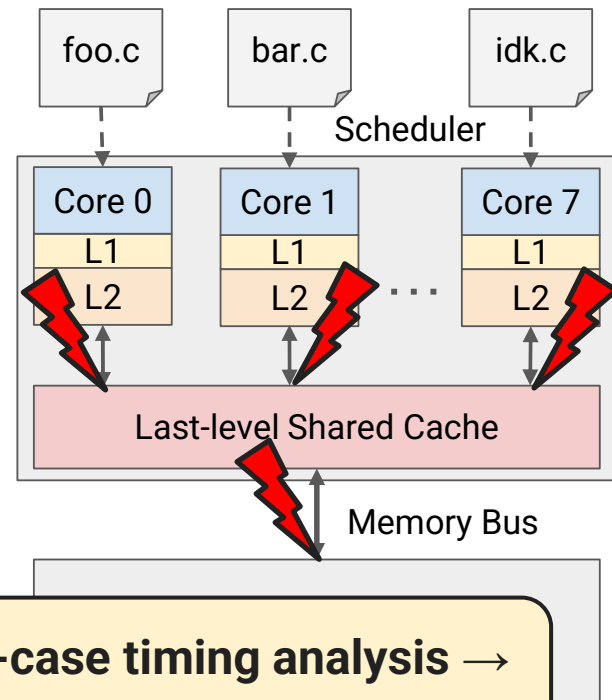

Challenges on Multicore

- Contention for shared resources can cause **interference** between tasks
- Example: **WCET slowdown** of PARSEC benchmarks due to interference



Challenges on Multicore

- Contention for shared resources can cause **interference** between tasks
- Example: **WCET slowdown** of PARSEC benchmarks due to interference



Potential interference → overly-conservative worst-case timing analysis → **over-provisioning** of hardware resources

Resource Contention: State of the Art

	Resource aware?	DAG support?	Resource control?	Dynamic control?	Timing analysis?	Takeaway
[Shi et al. RTAS '24] [Casini et al. RTAS '20] [Tessler et al. RTSS '23] [Zhao et al. RTNS '23]	✓	✓	✗	✗	✓	
CaM [Xu et al. RTAS '19] MMO [Sun et al. ECRTS '25]	✓	✗	✓	✗	✓	
DNA [Gifford et al. RTAS '20]	✓	✗	✓	✓	✗	

Resource Contention: State of the Art

	Resource aware?	DAG support?	Resource control?	Dynamic control?	Timing analysis?	Takeaway
[Shi et al. RTAS '24] [Casini et al. RTAS '20] [Tessler et al. RTSS '23] [Zhao et al. RTNS '23]	✓	✓	✗	✗	✓	No resource isolation, can still have interference
CaM [Xu et al. RTAS '19] MMO [Sun et al. ECRTS '25]	✓	✗	✓	✗	✓	
DNA [Gifford et al. RTAS '20]	✓	✗	✓	✓	✗	

Resource Contention: State of the Art

	Resource aware?	DAG support?	Resource control?	Dynamic control?	Timing analysis?	Takeaway
[Shi et al. RTAS '24] [Casini et al. RTAS '20] [Tessler et al. RTSS '23] [Zhao et al. RTNS '23]	✓	✓	✗	✗	✓	No resource isolation, can still have interference
CaM [Xu et al. RTAS '19] MMO [Sun et al. ECRTS '25]	✓	✗	✓	✗	✓	
DNA [Gifford et al. RTAS '20]	✓	✗	✓	✓	✗	

Resource Contention: State of the Art

	Resource aware?	DAG support?	Resource control?	Dynamic control?	Timing analysis?	Takeaway
[Shi et al. RTAS '24] [Casini et al. RTAS '20] [Tessler et al. RTSS '23] [Zhao et al. RTNS '23]	✓	✓	✗	✗	✓	No resource isolation, can still have interference
CaM [Xu et al. RTAS '19] MMO [Sun et al. ECRTS '25]	✓	✗	✓	✗	✓	Static allocation is resource inefficient
DNA [Gifford et al. RTAS '20]	✓	✗	✓	✓	✗	

Resource Contention: State of the Art

	Resource aware?	DAG support?	Resource control?	Dynamic control?	Timing analysis?	Takeaway
[Shi et al. RTAS '24] [Casini et al. RTAS '20] [Tessler et al. RTSS '23] [Zhao et al. RTNS '23]	✓	✓	✗	✗	✓	No resource isolation, can still have interference
CaM [Xu et al. RTAS '19] MMO [Sun et al. ECRTS '25]	✓	✗	✓	✗	✓	Static allocation is resource inefficient
DNA [Gifford et al. RTAS '20]	✓	✗	✓	✓	✗	

Resource Contention: State of the Art

	Resource aware?	DAG support?	Resource control?	Dynamic control?	Timing analysis?	Takeaway
[Shi et al. RTAS '24] [Casini et al. RTAS '20] [Tessler et al. RTSS '23] [Zhao et al. RTNS '23]	✓	✓	✗	✗	✓	No resource isolation, can still have interference
CaM [Xu et al. RTAS '19] MMO [Sun et al. ECRTS '25]	✓	✗	✓	✗	✓	Static allocation is resource inefficient
DNA [Gifford et al. RTAS '20]	✓	✗	✓	✓	✗	Dynamic allocation breaks timing guarantee

Rasco: Fine-grain Resource Control with Guarantees

	Resource aware?	DAG support?	Resource control?	Dynamic control?	Timing analysis?
Rasco [this work]	✓	✓	✓	✓	✓

We want the best of both worlds:

- ✓ Tight worst-case timing analysis via resource isolation
- ✓ Dynamic allocation of resources based on fine-grain needs

Research Questions

1. How do we design a task model that enables dynamic resource allocation and worst-case timing analysis?
2. Using this model, how do we allocate resources to improve the
 - resource efficiency,
 - average-case latency,
 - and hard real-time schedulabilityof DAG applications?


Contributions

- To answer RQ1, we propose a **resource-dependent multi-phase task model** which enables worst-case timing analysis under dynamic resource allocation
- To answer RQ2, we develop Rasco, a resource allocation and scheduling **co-design algorithm** for DAG applications on multicore
- We then implement a **prototype** of Rasco to evaluate the safety and utility of our approach in a real-time operating system

Talk Outline


1. Introduction/Background
- 2. The Resource-Dependent Multi-Phase Model**
3. Rasco: Resource Allocation and Scheduling Co-design
4. Numerical Evaluation
5. Prototype Evaluation and Overhead Accounting
6. Conclusion

Task Execution Phases




FFT with 10% of LLC and memory BW

Task Execution Phases

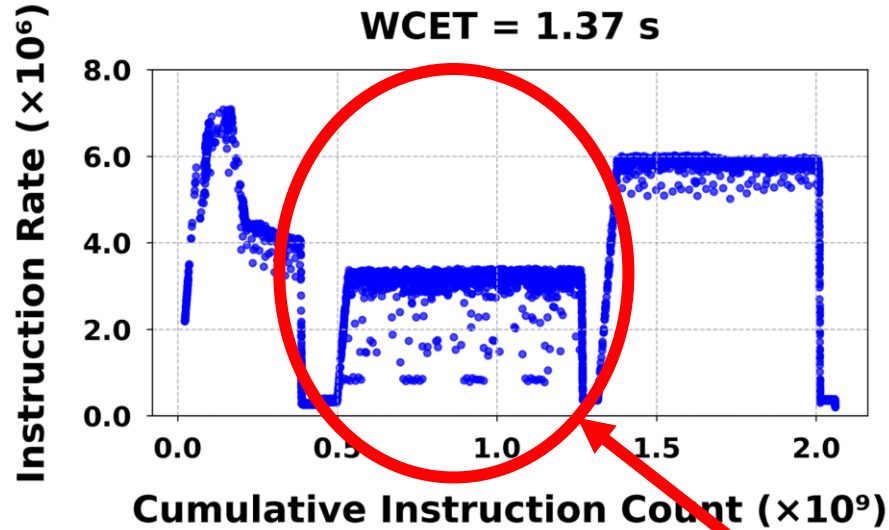


FFT with 10% of LLC and memory BW




FFT with 50% of LLC and memory BW

Task Execution Phases




FFT with 10% of LLC and memory BW




FFT with 50% of LLC and memory BW

Middle phase is highly resource intensive

Task Execution Phases



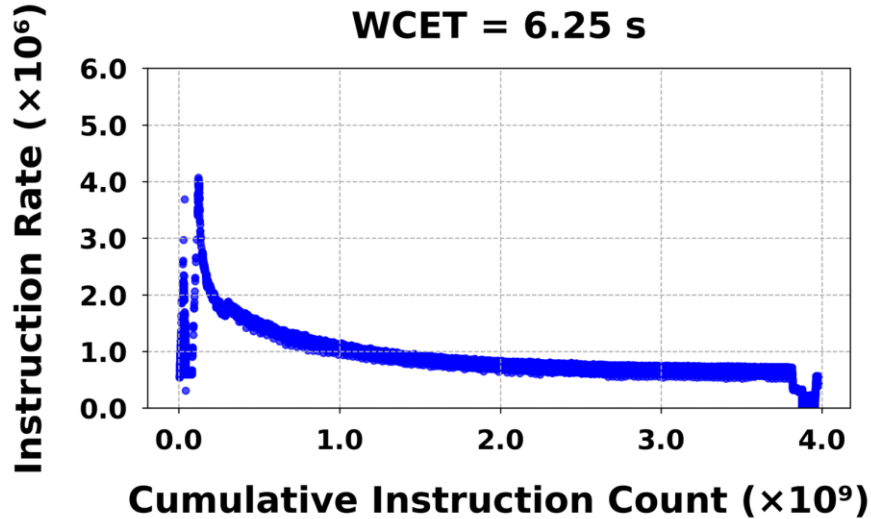
FFT with 10% of LLC and memory BW




FFT with 50% of LLC and memory BW

- Tasks have different execution phases
- Phases vary in resource intensity

A General Model for Resource-Dependent Phases

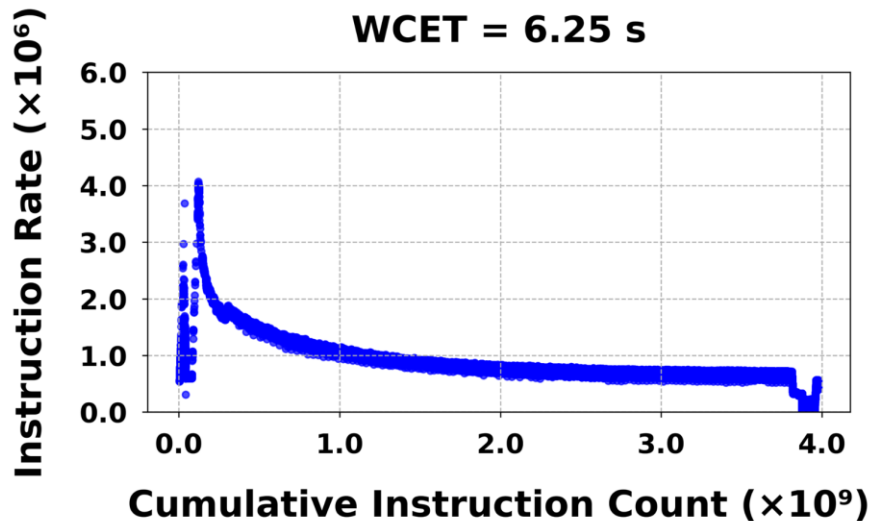


Canneal with 10% of LLC and mem BW




Canneal with 50% of LLC and mem BW

A General Model for Resource-Dependent Phases



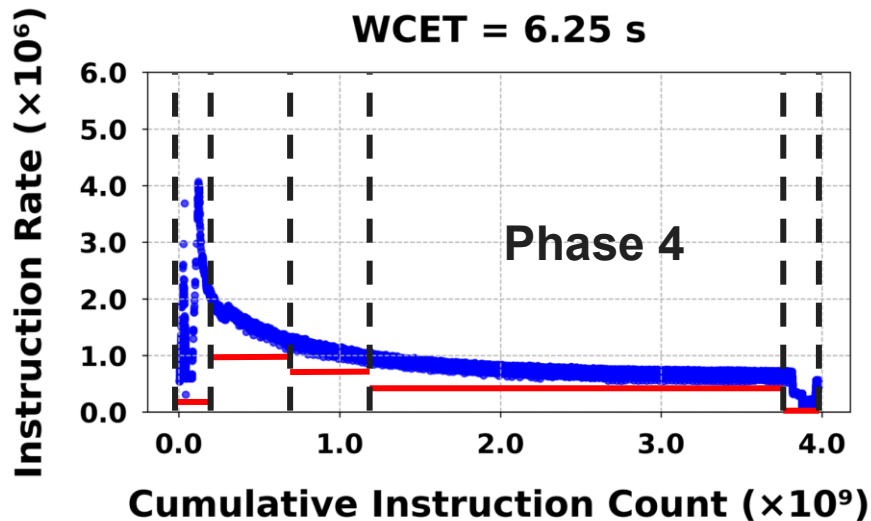
Canneal with 10% of LLC and mem BW




Canneal with 50% of LLC and mem BW

**Need a resource-dependent task model that is generalizable,
but still tight.**

A General Model for Resource-Dependent Phases




Canneal with 10% of LLC and mem BW




Canneal with 50% of LLC and mem BW

- Step 1: use changepoint detection to identify phases
- Step 2: compute worst-case instruction rates for each phase


WCET Under Dynamic Resource Allocation



WCET Under Dynamic Resource Allocation

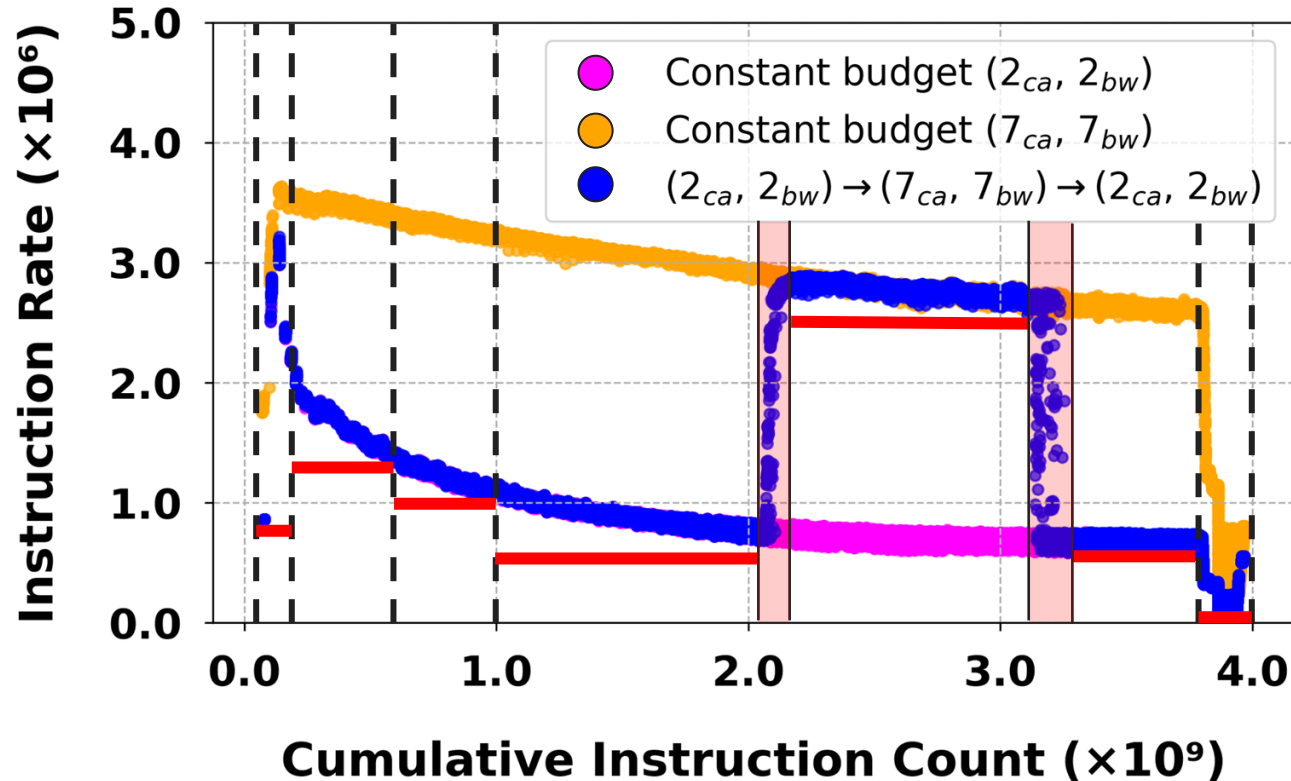


WCET Under Dynamic Resource Allocation




Delay between
resource
reconfiguration
and rate change

WCET Under Dynamic Resource Allocation




WCET Under Dynamic Resource Allocation





To get WCET under dynamic budget:

- Compose worst-case rates from constant budgets
- Incorporate delay at reconfiguration points

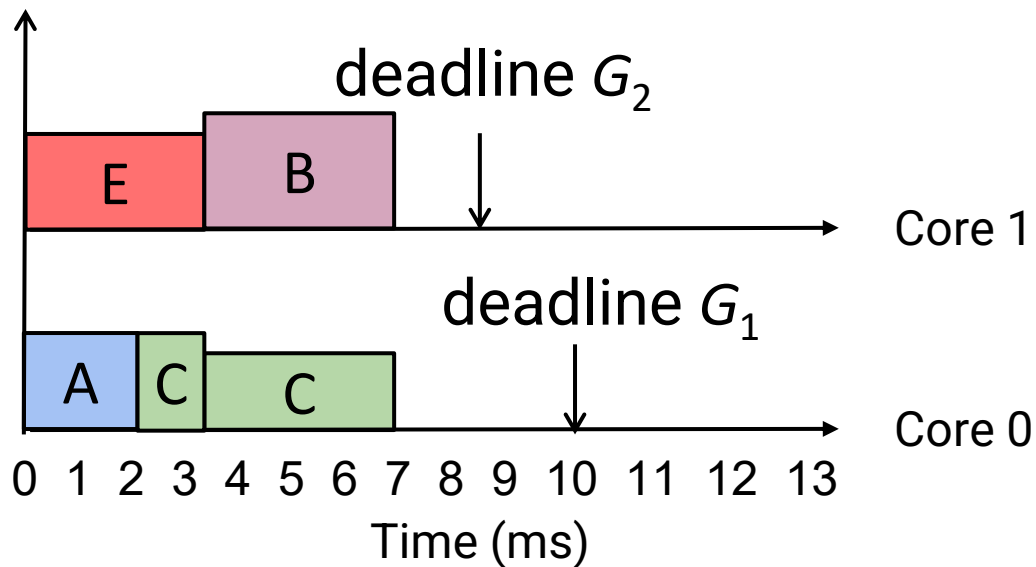
Research Questions

1. How do we design a task model that enables dynamic resource allocation and worst-case timing analysis? 
2. Using this model, how do we allocate resources to improve the
 - resource efficiency,
 - average-case latency,
 - and hard real-time schedulabilityof DAG applications?



Resource Allocation and Scheduling



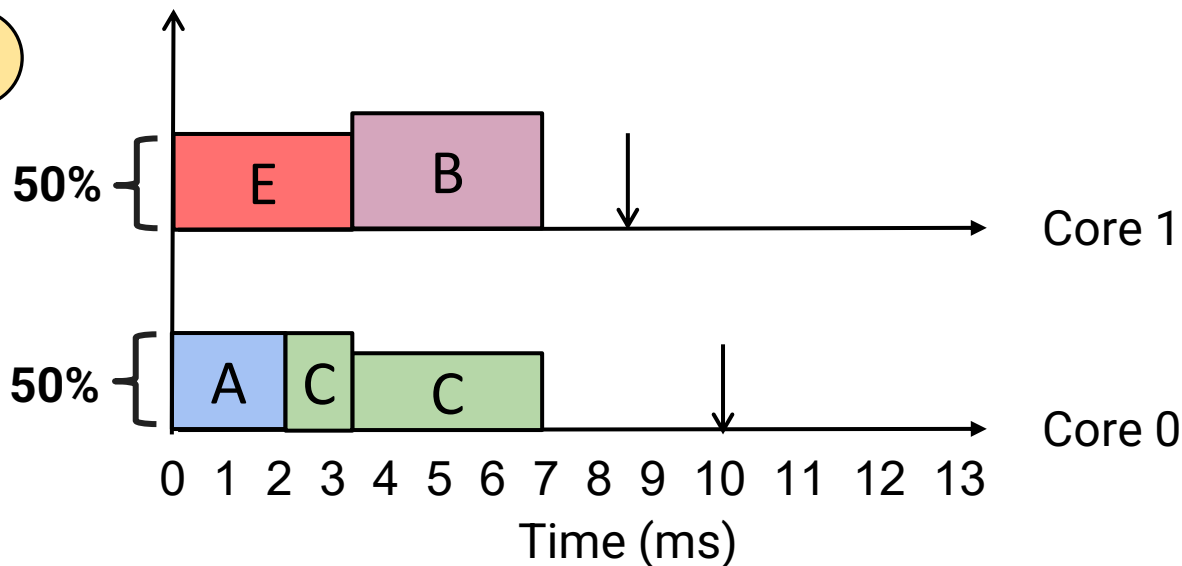
Suppose we are at time $t = 7$
and need to schedule D and F





Resource Allocation and Scheduling




Height of task = proportion of shared resources allocated





Resource Allocation and Scheduling



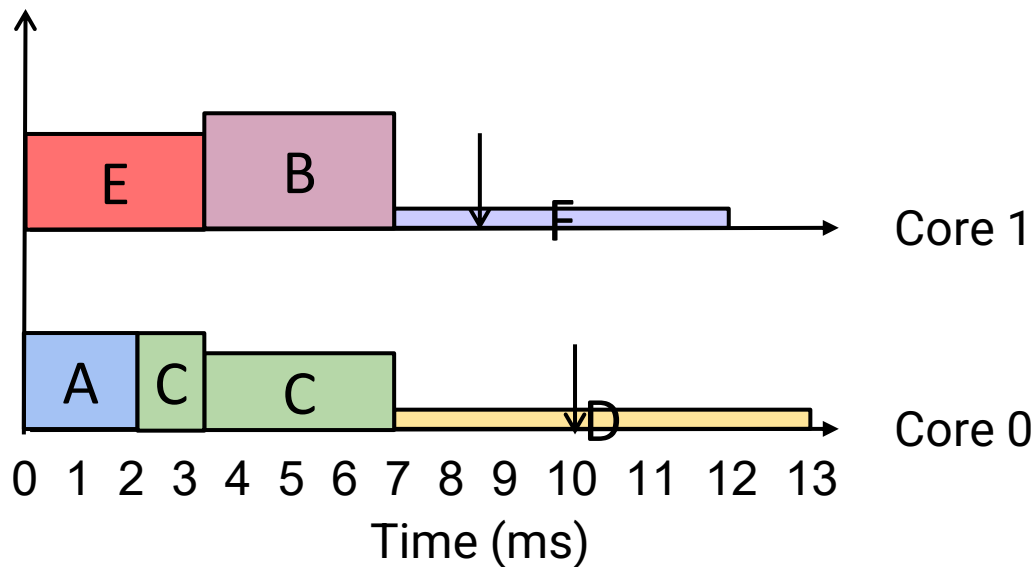

Height of task = proportion of shared resources allocated



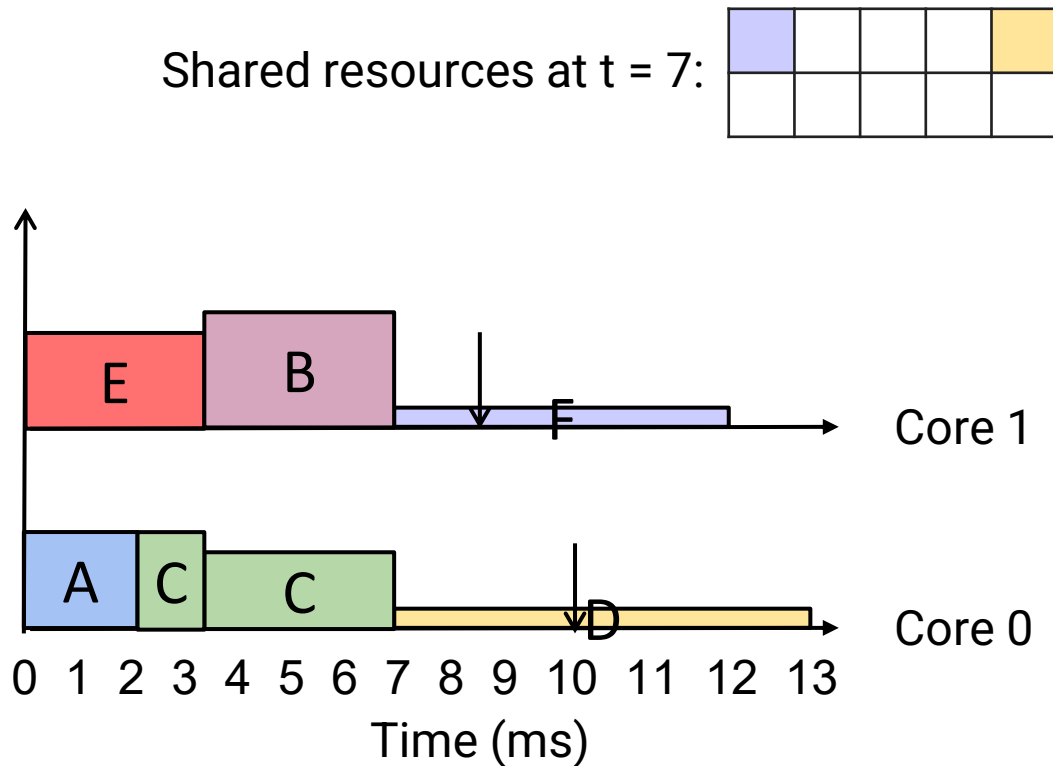


Resource Allocation and Scheduling



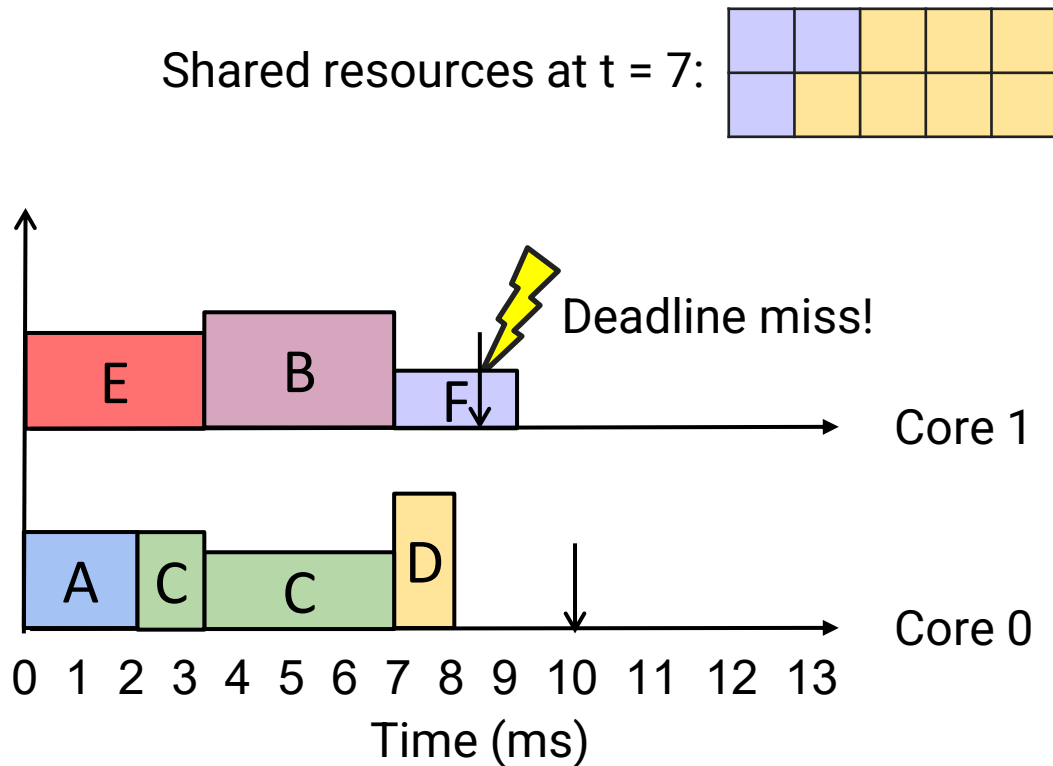


Shared resources at $t = 7$:



Resource Allocation and Scheduling



Resource Allocation and Scheduling




The Case for Co-Design

Maximizing resource efficiency \neq maximizing schedulability

The Case for Co-Design

Maximizing resource efficiency \neq maximizing schedulability




Resource allocation and scheduling must be co-designed.

Talk Outline


1. Introduction/Background
2. The Resource-Dependent Multi-Phase Model
- 3. Rasco: Resource Allocation and Scheduling Co-design**
4. Numerical Evaluation
5. Prototype Evaluation and Overhead Accounting
6. Conclusion

Rasco Algorithm Overview




Input: periodic DAG tasks and the multi-phase model for each task

Rasco Algorithm Overview




Input: periodic DAG tasks and the multi-phase model for each task







Output: task schedule with partition of shared resources at each scheduling point

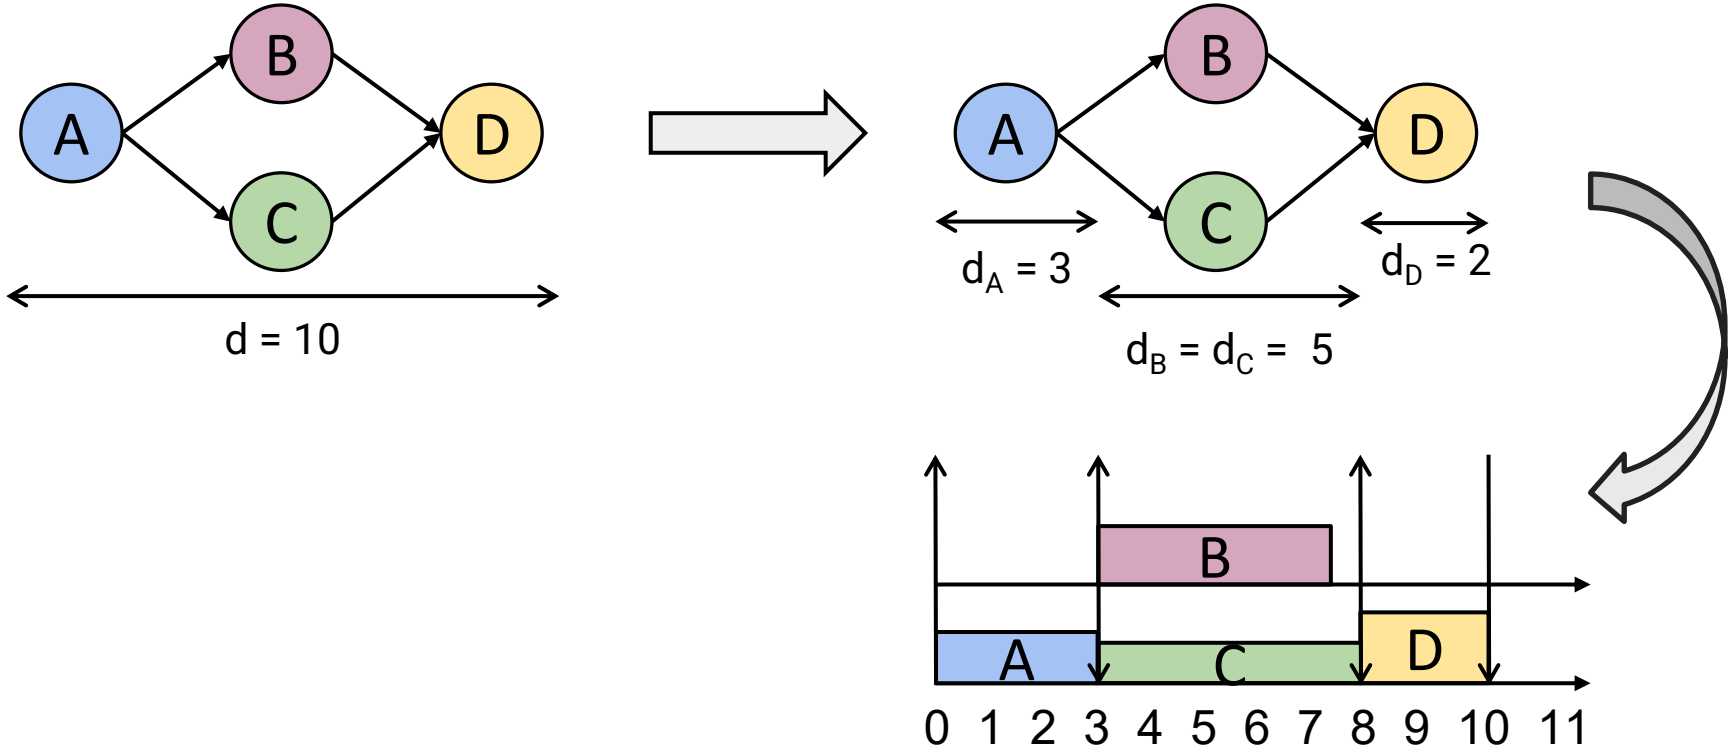
Rasco Pre-Processing: Deadline decomposition




Rasco Pre-Processing: Deadline decomposition




Rasco Pre-Processing: Deadline decomposition



Rasco Pre-Processing: Stack all DAG tasks (G_1, G_2)




Step 1: Start at $t = 0$, give out remaining resources at t

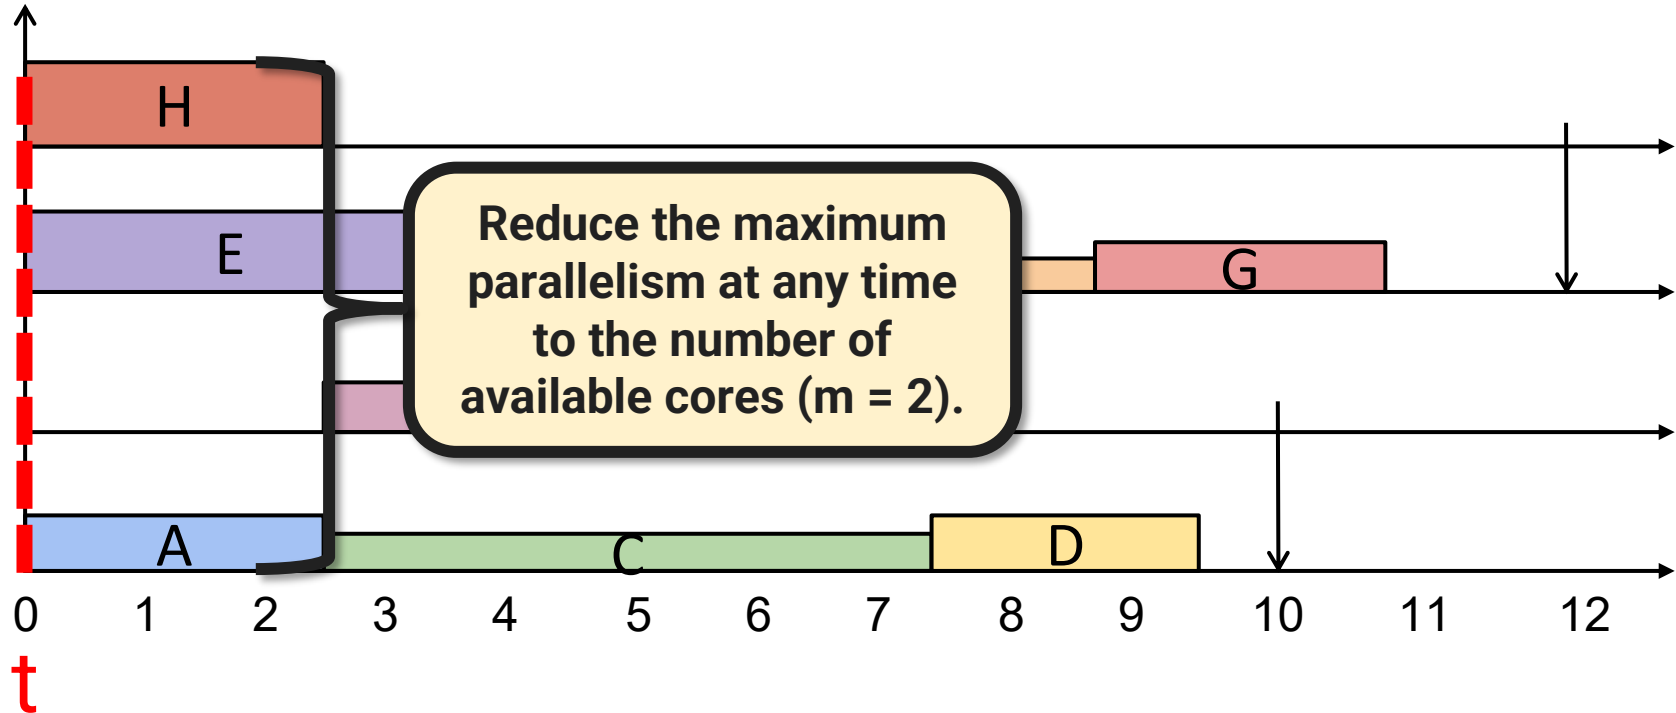


Step 1: Start at $t = 0$, give out remaining resources at t


Maximize resource efficiency at time t




Step 1: Start at $t = 0$, give out remaining resources at t




Step 2: Schedule $m = 2$ tasks with smallest deadlines



Step 2: Schedule $m = 2$ tasks with smallest deadlines




Step 3: Get next t, repeat




Step 3: Get next t, repeat

Maximize resource efficiency at t




Step 3: Get next t , repeat

Schedule m tasks with earliest deadlines




Step 3: Get next t , repeat




Step 3: Get next t, repeat

Maximize resource efficiency at t




Step 3: Get next t , repeat




Step 3: Get next t, repeat

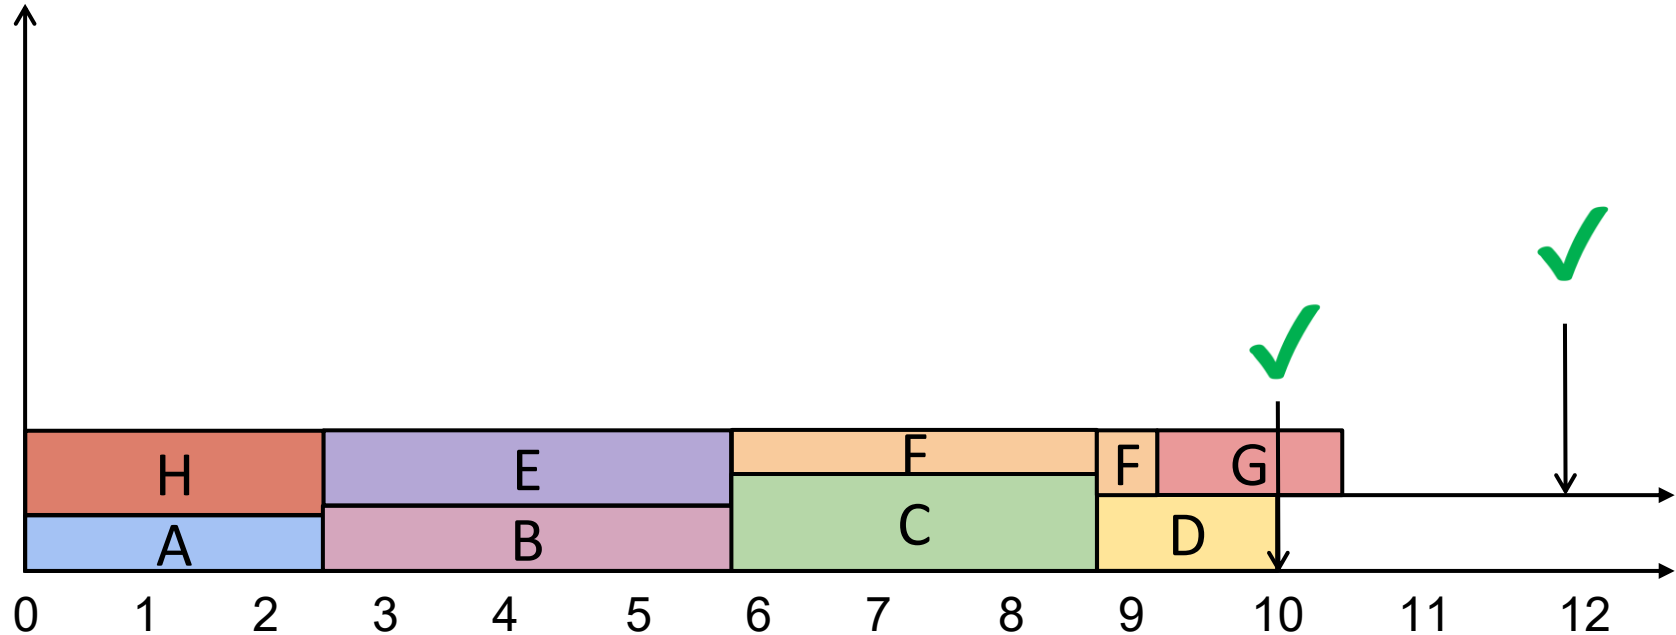
Maximize resource efficiency at t



Step 4: Finish and squash onto 2 cores



Step 5: Check if schedulable



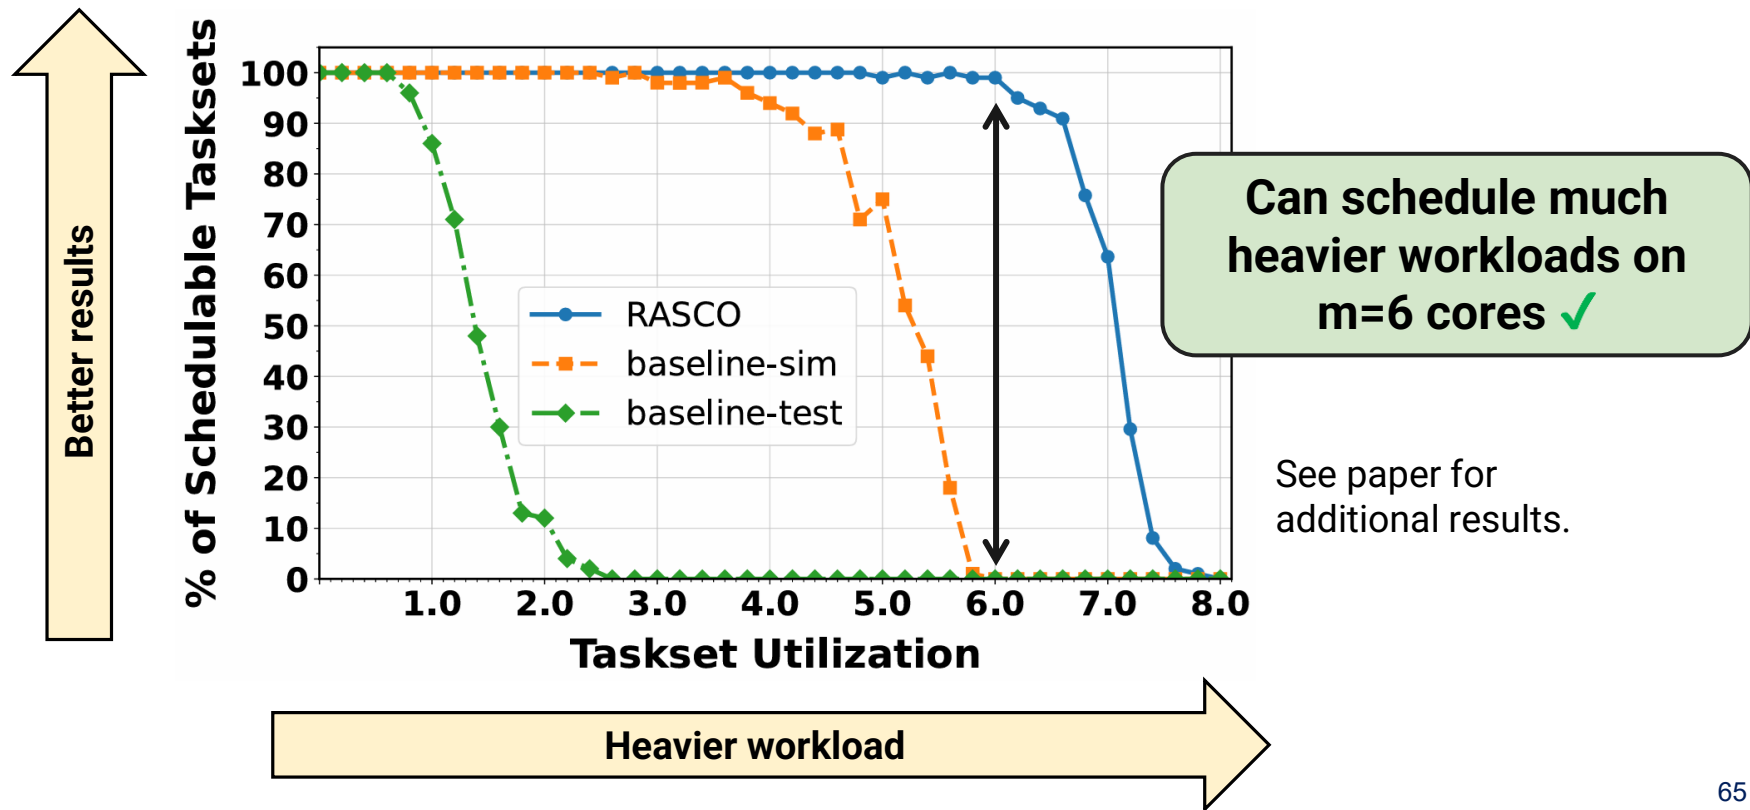
Talk Outline

1. Introduction/Background
2. The Resource-Dependent Multi-Phase Model
3. Rasco: Resource Allocation and Scheduling Co-design
- 4. Numerical Evaluation**
5. Prototype Evaluation and Overhead Accounting
6. Conclusion


Numerical Evaluation Setup

- Profiled benchmarks from PARSEC and SPLASH2x
- Used Intel's CAT and MemGuard to partition shared resources
- Constructed multi-phase models using changepoint detection
- Randomly generated 100 tasksets per utilization step [X. Dai. [dag-gen-rnd](#)]
- Ran Rasco on each taskset
- Compared **schedulability** and **latency** against a state-of-the-art deadline decomposition method using even partition of resources to cores
 - Schedulability test: **baseline-test**
 - Simulated schedule under global EDF: **baseline-sim**

Schedulability Results




Latency Results



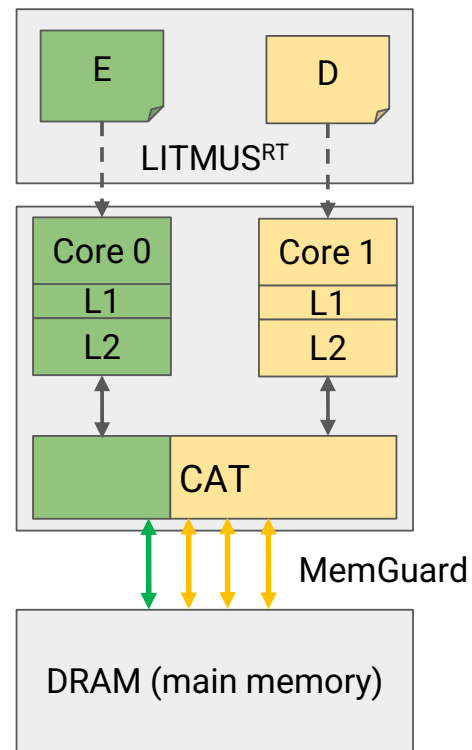
Talk Outline

1. Introduction/Background
2. The Resource-Dependent Multi-Phase Model
3. Rasco: Resource Allocation and Scheduling Co-design
4. Numerical Evaluation
- 5. Prototype Evaluation and Overhead Accounting**
6. Conclusion

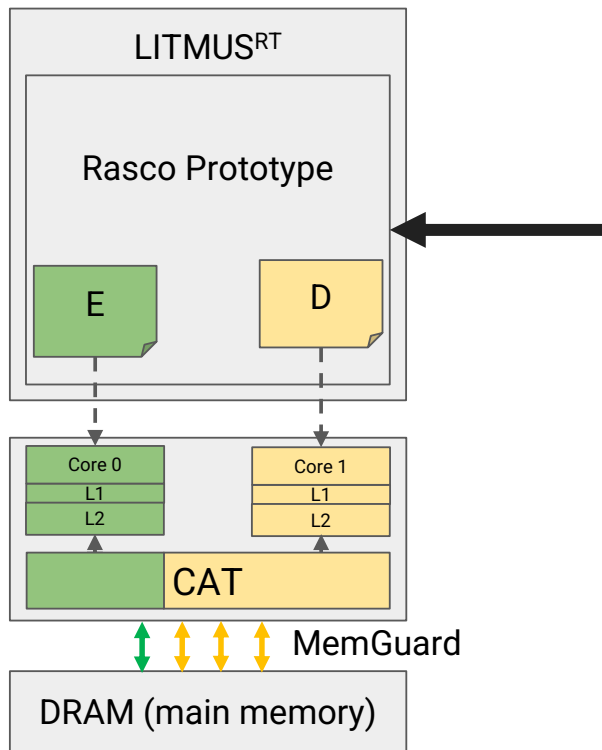
Prototype in LITMUS^{RT}



Prototype state at t_2




Rasco Prototype Features



1. Time-triggered table-driven scheduling in LITMUS^{RT}
2. Runtime reconfiguration of CAT + MemGuard in LITMUS^{RT}
3. Logic for handling job under-runs
4. Early releasing and work stealing

Overhead-Aware Rasco Extension


	Min (μ s)	Mean (μ s)	99 th (μ s)	Max (μ s)
Rasco Scheduling	0.02	0.03	0.05	18.10
CAT + MemGuard	1.66	2.53	5.80	23.30



Observed small runtime overheads ✓

Overhead-Aware Rasco Extension


	Min (μ s)	Mean (μ s)	99 th (μ s)	Max (μ s)
Rasco Scheduling	0.02	0.03	0.05	18.10
CAT + MemGuard	1.66	2.53	5.80	23.30



99.6% of the tasksets had less than 2% overhead ✓


Empirical Evaluation on Prototype

- Ran Rasco's output schedules on our prototype



Empirical Evaluation on Prototype


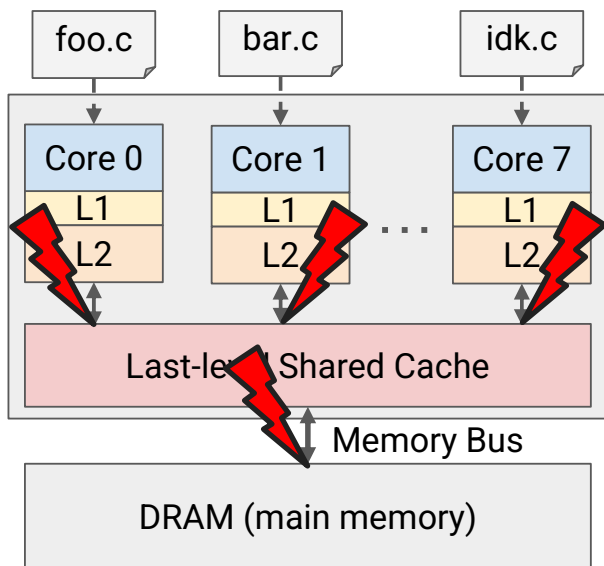
- Ran Rasco's output schedules on our prototype



Empirical schedulability always exceeds theoretical guarantee ✓

Recap

- The move to multicore introduces challenges for timing analysis
- Resource contention → overly-conservative analysis → over-provisioning



Recap

- No prior work had achieved the **resource efficiency** of fine-grained **dynamic resource allocation** while providing **hard real-time guarantees**

	Resource aware?	DAG support?	Resource control?	Dynamic control?	Timing analysis?
Rasco	✓	✓	✓	✓	✓

Conclusion

- Proposed a **resource-dependent multi-phase model** which enables worst-case timing analysis under dynamic resource allocation
- Developed a resource allocation and scheduling **co-design algorithm** for DAG applications on multicore that improves
 - resource efficiency,
 - latency, and
 - schedulability
- Implemented a **prototype** of Rasco to evaluate the safety and utility of our approach in a real-time operating system

Conclusion

- Proposed a **resource-dependent multi-phase model** which enables worst-case timing analysis under dynamic resource allocation
- Developed a resource allocation and scheduling **co-design algorithm** for DAG applications on multicore that improves
 - resource efficiency,
 - latency, and
 - schedulability
- Implemented a **prototype** of Rasco to evaluate the safety and utility of our approach in a real-time operating system

Thank you!

<https://github.com/abbyeisenklam/Rasco>